


Spark Introduction

@CrazyJvm

什么是Spark？

Apache Spark is an open source cluster computing system that aims to make data analytics fast --both fast to run and fast to write.

就这么简单！



Driver程序启动很多workers, 然后workers在(分布式)文件系统中读取数据后转化为RDD, 最后对RDD在内存中进行缓存和计算。

为什么Spark的性能高？

- 1、in-memory computation
- 2、general computation graphs

可参考 [Delay scheduling : A simple technique for achieving locality and fairness in cluster scheduling](#)

提供了哪些语言的API？

Scala

Java

Python

怎样运行Spark？

- Local
- Standalone
- Mesos
- YARN

可参考：@明风Andy 出品的 [基于Spark on Yarn的淘宝数据挖掘平台](#)

Scala光速入门

玩Spark最好会点Scala。

直接引用Matei-Zaharia之前介绍Scala的示例

Scala语言？

- 1、基于JVM的FP+OO
- 2、静态类型
- 3、和JAVA可以互操作

可以在Scala的命令行里试试！

Scala--变量声明

```
var x : Int = 7
```

```
val x = 7 //自动类型推导
```

```
val y = "hi" //只读的 相当于Java里的final
```

Scala--函数

```
def square(x : Int) : Int = x * x
```

```
def square(x : Int) : Int = {
```

```
    x * x //在block中的最后一个值将被返回
```

```
}
```

```
def announce(text : String) {
```

```
    println(text)
```

```
}
```

Scala--泛型

```
var arr = new Array[Int](8)
```

```
var lst = List(1,2,3) //1st的类型是List[Int]
```

```
//索引访问
```

```
arr(5) = 7
```

```
println(lst(1))
```

Scala--FP的方式处理集合

```
val list = List(1,2,3)
```

```
list.foreach(x => println(x)) // 打印出1,2,3
```

```
list.foreach(println)
```

```
list.map(x => x + 2) // List(3,4,5)
```

```
list.map(_ + 2)
```

```
list.filter(x => x % 2 == 1) // List(1,3)
```

```
list.filter(_ % 2 == 1)
```

```
list.reduce((x,y) => x + y) // 6
```

```
list.reduce(_ + _)
```

Scala--闭包

```
(x : Int) => x + 2
```

```
x => x + 2
```

```
_ + 2
```

```
x => {
```

```
    val numberToAdd = 2
```

```
    x + numberToAdd
```

```
}
```

//如果闭包很长, 可以考虑作为参数传入

```
def addTwo(x : Int) : Int = x + 2
```

```
list.map(addTwo)
```

回到Spark

Spark : 从RDD说起

RDDs : resilient distributed datasets

1. **immutable** collections of objects spread across a cluster
2. build through parallel transformations(map,filter,groupBy,join etc)
3. automatically rebuild on failure (基于lineage和checkpoint)
4. different storage level (可(序列化)存在内存或者硬盘中)

单独看下StorageLevel

```
val NONE = new StorageLevel(false, false, false)
val DISK_ONLY = new StorageLevel(true, false, false)
val DISK_ONLY_2 = new StorageLevel(true, false, false, 2)
val MEMORY_ONLY = new StorageLevel(false, true, true)
val MEMORY_ONLY_2 = new StorageLevel(false, true, true, 2)
val MEMORY_ONLY_SER = new StorageLevel(false, true, false)
val MEMORY_ONLY_SER_2 = new StorageLevel(false, true, false, 2)
val MEMORY_AND_DISK = new StorageLevel(true, true, true)
val MEMORY_AND_DISK_2 = new StorageLevel(true, true, true, 2)
val MEMORY_AND_DISK_SER = new StorageLevel(true, true, false)
val MEMORY_AND_DISK_SER_2 = new StorageLevel(true, true, false, 2)
```


来建立几个RDD

sc就是SparkContext,可以看做用户在spark上的入口。

```
scala> val text = sc.textFile("README.md")
13/12/17 14:00:57 INFO storage.MemoryStore: ensureFreeSpace(44405) called with curMem=67016, maxMem=340147568
13/12/17 14:00:57 INFO storage.MemoryStore: Block broadcast_2 stored as values to memory (estimated size 43.4 KB, free 324.3 MB)
text: org.apache.spark.rdd.RDD[String] = MappedRDD[5] at textFile at <console>:12

scala> val logon = sc.textFile("hdfs://server1:9000/data/logon/logon-1.log")
13/12/17 14:01:00 INFO storage.MemoryStore: ensureFreeSpace(44365) called with curMem=111421, maxMem=340147568
13/12/17 14:01:00 INFO storage.MemoryStore: Block broadcast_3 stored as values to memory (estimated size 43.3 KB, free 324.2 MB)
logon: org.apache.spark.rdd.RDD[String] = MappedRDD[7] at textFile at <console>:12
```

RDD到底是啥玩意儿？

internally, each RDD is characterized by five main properties:

- * - A list of partitions
- * - A function for computing each split
- * - A list of dependencies on other RDDs
- * - Optionally, a Partitioner for key-value RDDs (e.g. to say that the RDD is hash-partitioned)
- * - Optionally, a list of preferred locations to compute each split on (e.g. block locations for an HDFS file)

进一步看下RDD

- RDDs (可以从一系列数据源创建)
- Transformations (延迟执行)
- Action (def runJob[T, U: ClassManifest]
(rdd: RDD[T],
func: Iterator[T] => U): Array[U])

Transformation and Action

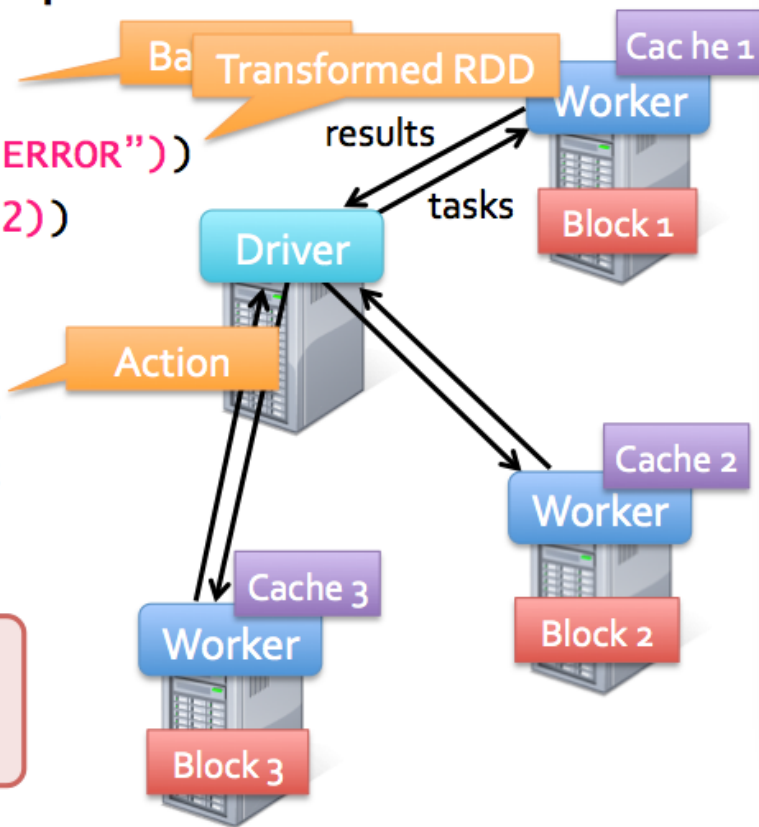
Transformations	<ul style="list-style-type: none"><code>map(f : T ⇒ U)</code> : <code>RDD[T] ⇒ RDD[U]</code><code>filter(f : T ⇒ Bool)</code> : <code>RDD[T] ⇒ RDD[T]</code><code>flatMap(f : T ⇒ Seq[U])</code> : <code>RDD[T] ⇒ RDD[U]</code><code>sample(fraction : Float)</code> : <code>RDD[T] ⇒ RDD[T]</code> (Deterministic sampling)<code>groupByKey()</code> : <code>RDD[(K, V)] ⇒ RDD[(K, Seq[V])]</code><code>reduceByKey(f : (V, V) ⇒ V)</code> : <code>RDD[(K, V)] ⇒ RDD[(K, V)]</code><code>union()</code> : <code>(RDD[T], RDD[T]) ⇒ RDD[T]</code><code>join()</code> : <code>(RDD[(K, V)], RDD[(K, W)]) ⇒ RDD[(K, (V, W))]</code><code>cogroup()</code> : <code>(RDD[(K, V)], RDD[(K, W)]) ⇒ RDD[(K, (Seq[V], Seq[W]))]</code><code>crossProduct()</code> : <code>(RDD[T], RDD[U]) ⇒ RDD[(T, U)]</code><code>mapValues(f : V ⇒ W)</code> : <code>RDD[(K, V)] ⇒ RDD[(K, W)]</code> (Preserves partitioning)<code>sort(c : Comparator[K])</code> : <code>RDD[(K, V)] ⇒ RDD[(K, V)]</code><code>partitionBy(p : Partitioner[K])</code> : <code>RDD[(K, V)] ⇒ RDD[(K, V)]</code>
Actions	<ul style="list-style-type: none"><code>count()</code> : <code>RDD[T] ⇒ Long</code><code>collect()</code> : <code>RDD[T] ⇒ Seq[T]</code><code>reduce(f : (T, T) ⇒ T)</code> : <code>RDD[T] ⇒ T</code><code>lookup(k : K)</code> : <code>RDD[(K, V)] ⇒ Seq[V]</code> (On hash/range partitioned RDDs)<code>save(path : String)</code> : Outputs RDD to a storage system, e.g., HDFS

举个经典的例子

```
val lines = spark.textFile("hdfs://...")  
val errors = lines.filter(_.startsWith("ERROR"))  
val messages = errors.map(_.split('\t')(2))  
messages.cache()
```

```
messages.filter(_.contains("foo")).count  
messages.filter(_.contains("bar")).count  
...
```

Result: scaled to 1 TB data in 5-7 sec
(vs 170 sec for on-disk data)



可以略微简化下代码

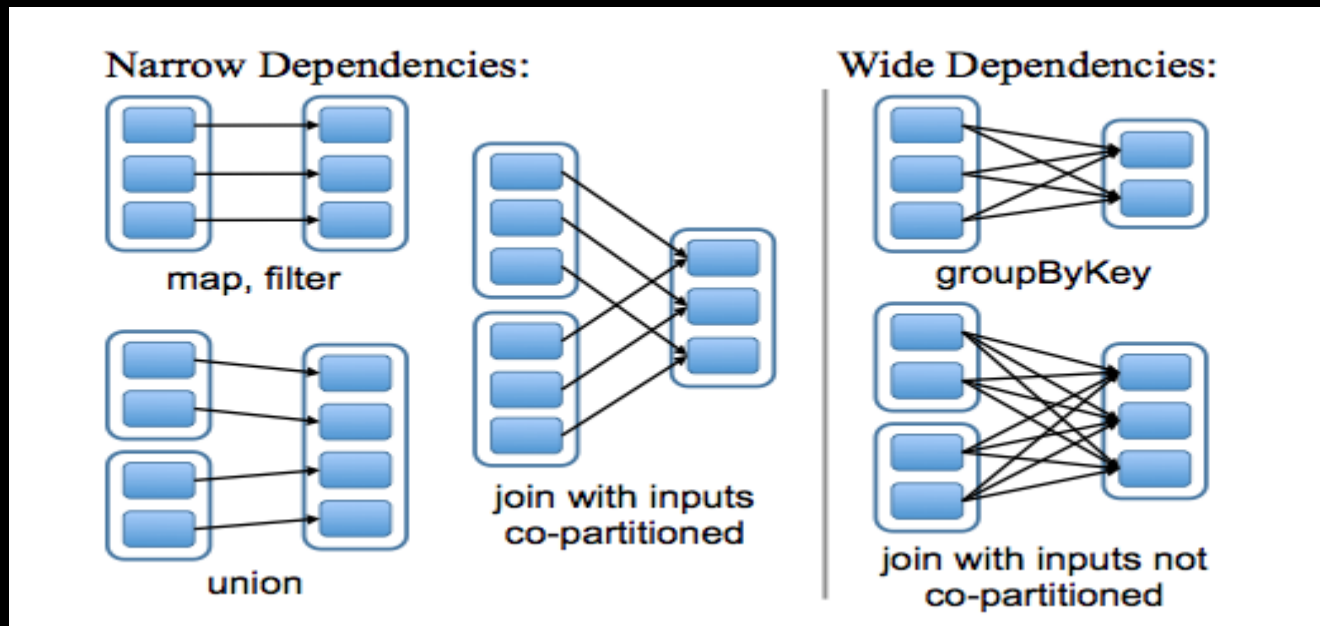
```
val message=spark.textFile("hdfs://...").filter(_.startsWith("ERROR")).map(_._split("\t")(2)).cache  
def getStatistic(keyword) : Long = message.filter(_._contains(keyword)).count
```

关于Cache多说几句

目前RDD的storage level一旦设定就不能再改变。以后可能会允许改变。(当然你可以使用unpersist去删除cache)

```
/**
 * Set this RDD's storage level to persist its values across operations after the first time
 * it is computed. This can only be used to assign a new storage level if the RDD does not
 * have a storage level set yet..
 */
def persist(newLevel: StorageLevel): RDD[T] = {
  // TODO: Handle changes of StorageLevel
  if (storageLevel != StorageLevel.NONE && newLevel != storageLevel) {
    throw new UnsupportedOperationException(
      "Cannot change storage level of an RDD after it was already assigned a level")
  }
  storageLevel = newLevel
  // Register the RDD with the SparkContext
  sc.persistentRdds(id) = this
  this
}
```

容错



基于lineage(子RDD挂了可以从父RDD再次计算得来)和checkpoint(物化)

接下来，举几个RDD的例子。

RDD的创建

1、直接从集合转化

```
sc.parallelize(List(1,2,3,4,5,6,7,8,9,10))
```

2、从各种(分布式)文件系统来

```
sc.textFile("README.md")
```

```
sc.textFile("hdfs://xxx")
```

3、从现存的任何Hadoop InputFormat而来

```
sc.hadoopFile(keyClass, valClass, inputFormat, conf)
```

Transformation

```
val nums = sc.parallelize(List(1,2,3,4,5,6,7,8,9))
```

```
nums: org.apache.spark.rdd.RDD[Int] = ParallelCollectionRDD[0] at parallelize at <console>:  
12
```

```
val squares = nums.map(x => x*x)
```

```
squares: org.apache.spark.rdd.RDD[Int] = MappedRDD[1] at map at <console>:14
```

```
val even = nums.filter(_ % 2 == 0)
```

```
even: org.apache.spark.rdd.RDD[Int] = FilteredRDD[2] at filter at <console>:14
```

Action

squares collect

res0: Array[Int] = Array(1, 4, 9, 16, 25, 36, 49, 64, 81)

even collect

res1: Array[Int] = Array(2, 4, 6, 8)

nums reduce (_ + _)

45

nums take 5

Array(1, 2, 3, 4, 5)

nums count

9

K-V style RDD

```
val rdd = sc.parallelize(List(("A",1), ("B",2), ("C",3), ("A",4), ("B",5)))
```

```
val rbk = rdd.reduceByKey(_+_).collect  
Array((A,5), (B,7), (C,3))
```

```
val gbk = rdd.groupByKey.collect  
Array((A,ArrayBuffer(1, 4)), (B,ArrayBuffer(2, 5)), (C,ArrayBuffer(3)))
```

```
val sbk = rdd.sortByKey().collect //注意这里sortByKey的小括号不能省。  
Array((A,1), (A,4), (B,2), (B,5), (C,3))
```

K-V style RDD 续

```
val player = sc.parallelize(List(("ACMILAN","KAKA"),("ACMILAN","BT"),("GUANGZHOU","ZHENGZHI")))
```

```
val team = sc.parallelize(List(("ACMILAN",5),("GUANGZHOU",3)))
```

```
player.join(team)
```

```
Array((GUANGZHOU,(ZHENGZHI,3)), (ACMILAN,(KAKA,5)), (ACMILAN,(BT,5)))
```

```
player.cogroup(team)
```

```
Array((GUANGZHOU,(ArrayBuffer(ZHENGZHI),ArrayBuffer(3))), (ACMILAN,(ArrayBuffer(KAKA, BT),  
ArrayBuffer(5))))
```

注意怎样控制reduce task的数量:

```
xx.reduceByKey(_+_ ,10)    xx.groupByKey(5)
```

会写wordcount程序了吗？

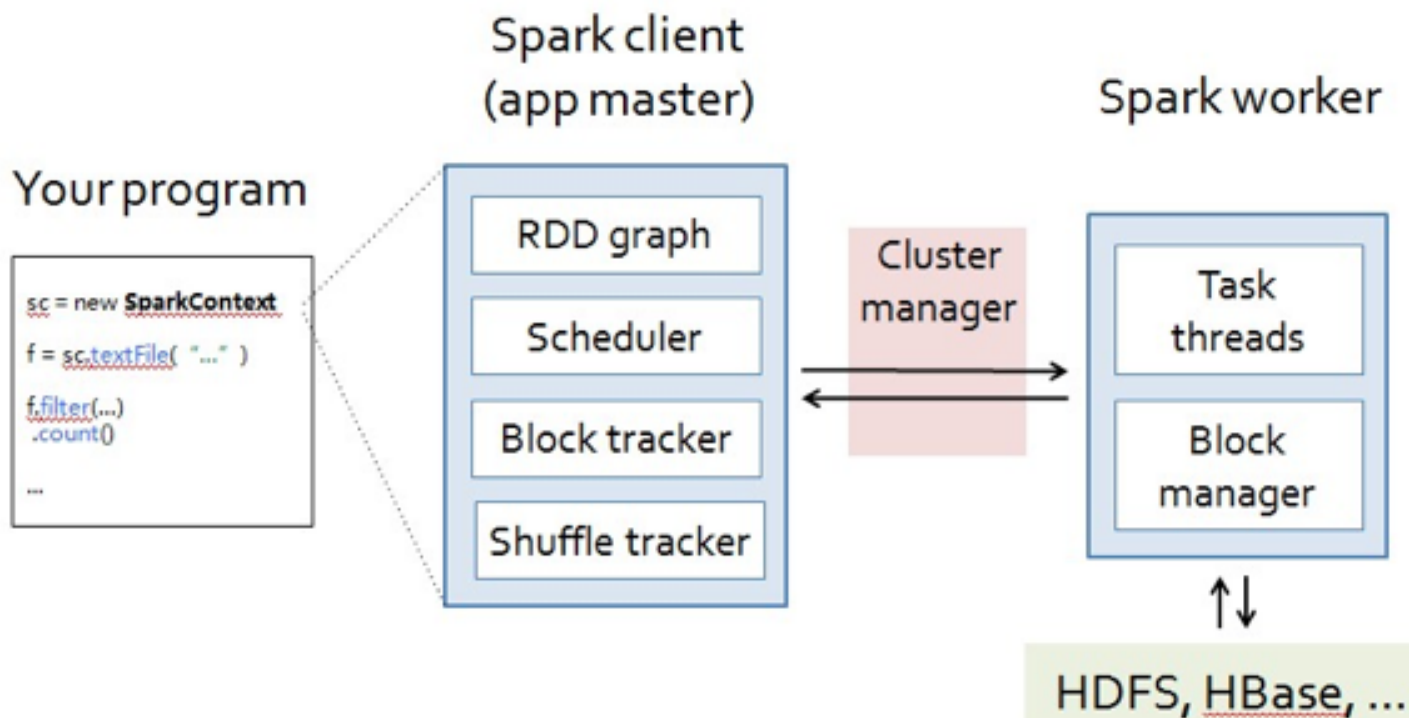
word count

```
sc.textFile("README.md").flatMap(_.split(' ')).map((_,1)).reduceByKey(_+_)
```

感受一下 r(^ ^)h


接下来看下作业调度相关内容

Spark组件



RDD Graph

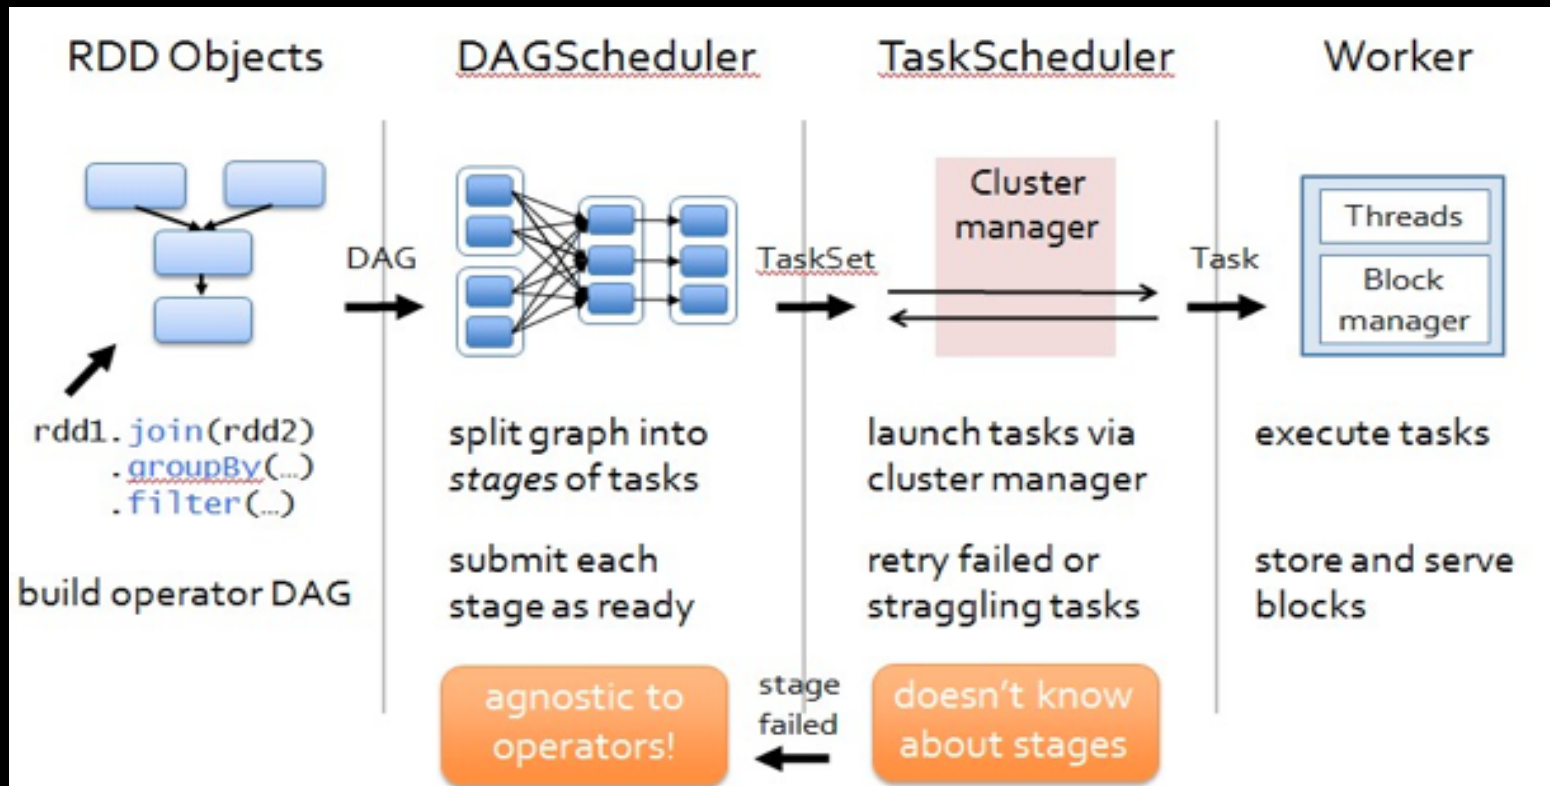
从Partition级别来看，每一个partition都会分配一个Task



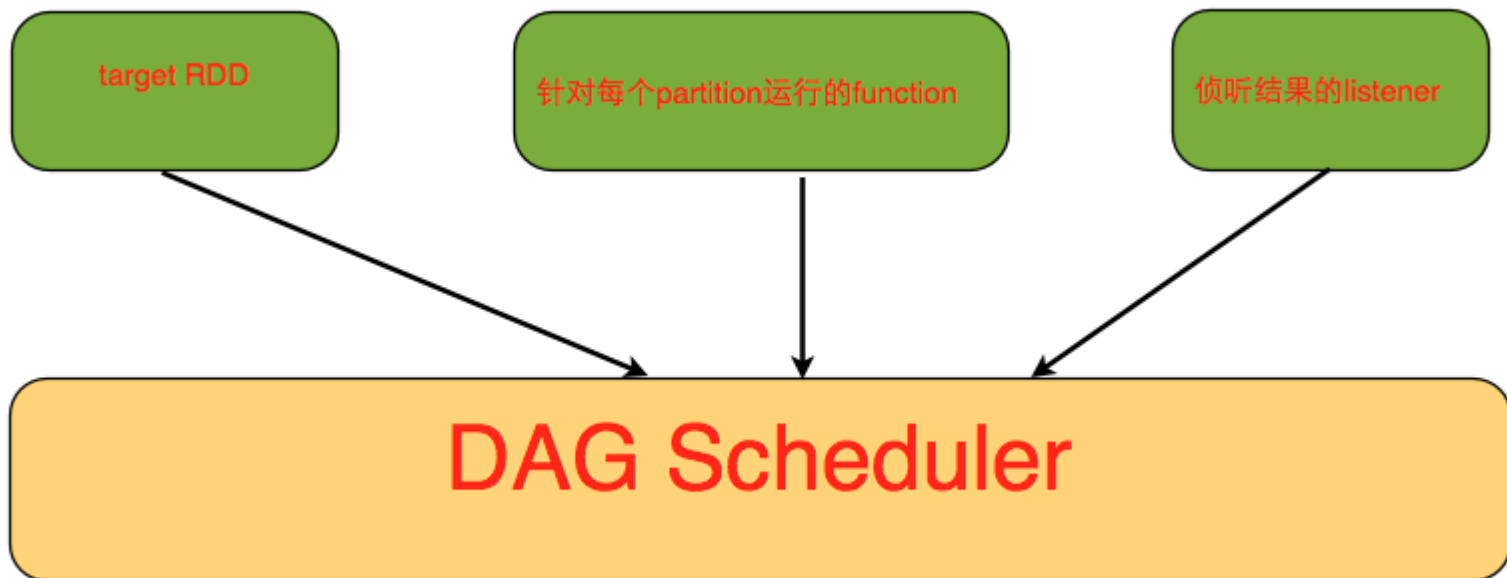
数据本地性

- 1、文件系统的本地性
- 2、内存本地性
- 3、内存假如被evict掉后退化为文件系统的本地性

作业与任务调度



DAAG Scheduler



DAG Scheduler的角色

- 1、为每个Job分割stage, 同时会决定最佳路径, 并且DAG Scheduler会记录哪个RDD或者stage的输出被物化, 从而来找到一个最优调度方案。
- 2、将TaskSet传给TaskScheduler
- 3、重新提交那些输出lost的stage

DAG Scheduler优化


1、某个stage内pipeline执行

对比： $1+1+1+1+1=5$ (1)

$1+1=2$ $2+1=3$ $3+1=4$ $4+1=5$ (2)

2、基于partition选择合适的join算法最小化shuffle


3、重用已经cache过的数据



TaskScheduler

- 1、提交tasks到集群并执行，假如出错就重试。
- 2、假如shuffle输出lost就报告fetch failed错误
- 3、遇到straggle task需要放到别的node上重试

实现：

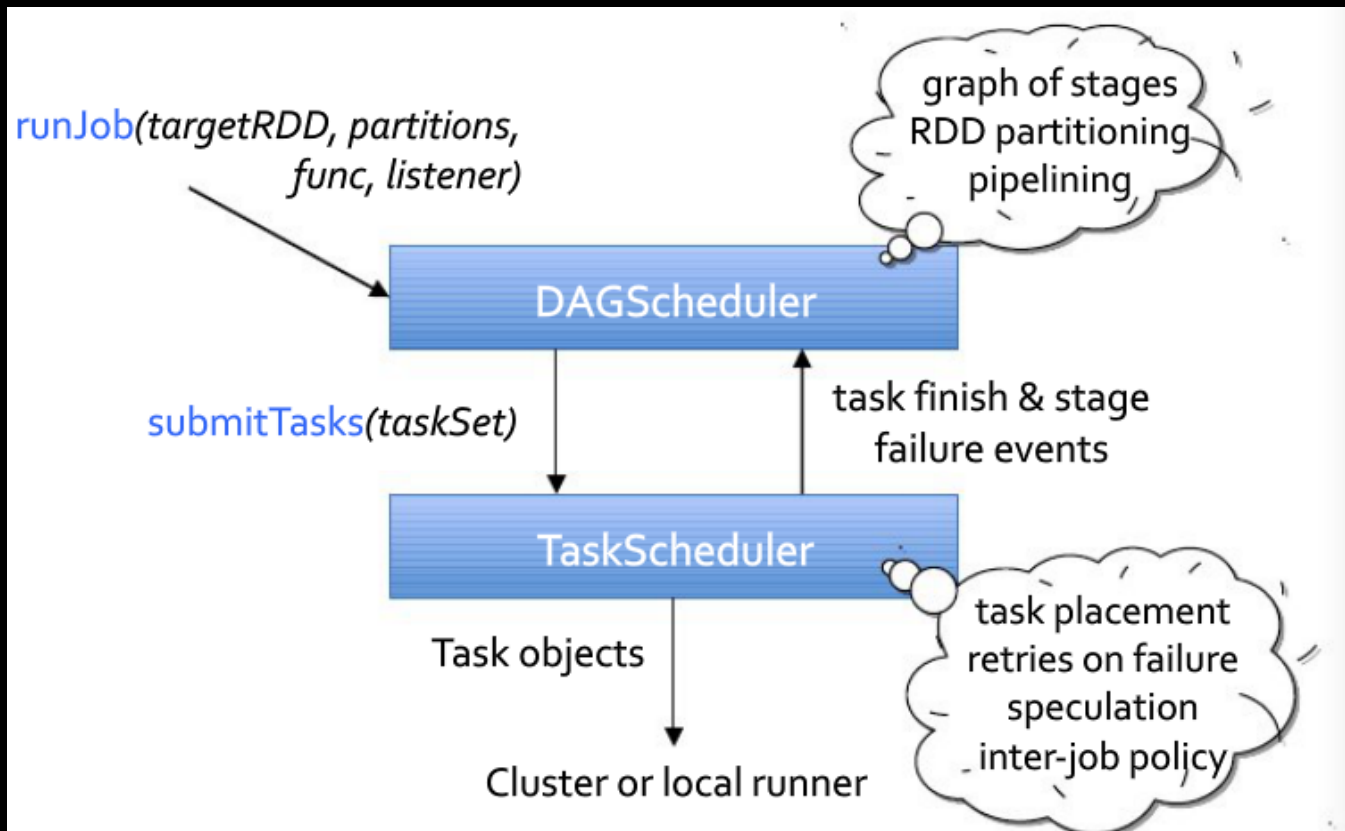


Task细节



物化Shuffle输出到内存或者硬盘以便重试

Job调度执行流程



Broadcast variables

BT style的变量广播

场景:lookup表, mapside join

参考:[Performance and Scalability of Broadcast in Spark](#)

Accumulators

类似于MapReduce中的counter

最后谈下Spark(性能)的优化

CASE 1

问题：序列化Task太大

解决：大对象使用broadcast variables

ps: spark 0.9版本中task序列化大小大于100k将会打印警告日志

参考：<https://spark-project.atlassian.net/browse/SPARK-590>

CASE 2

问题: `val rdd = xx.map(..).filter(..).filter(..)` ①
`rdd.map(xx).reduceBy...` ②

经过语句①, ②可能会产生很多小任务

解决: 使用`coalesce`或者`repartition`

CASE 3

问题：每个record的开销较大

典型场景 `rdd.map{x=>conn=getDBConn;conn.write(x.toString);conn.close}`

解决方案：

`rdd.mapPartitions(records=>conn=getDBConn;for(item <- records) write(item.toString);conn.close)`

CASE 4

问题:任务执行速度倾斜

解决:

1、数据倾斜(一般是partition key取的不好)

考虑其它的并行处理方式

中间可以加入一步aggregation

2、Worker倾斜(在某些worker上的executor不给力)

设置spark.speculation=true

把那些持续不给力的node去掉

CASE 5

问题: 不设置spark.local.dir

这是spark写shuffle输出的地方

解决: 可以设置一组disk

spark.local.dir=/mnt1/spark, /mnt2/spark, /mnt3/spark

CASE 6

问题: reducer数量不合适

解决: 看实际情况(请忽略这句废话!)

太多的reducer, 造成很多的小任务, 以此产生很多启动任务的开销。

太少的reducer, 慢! 并且可能造成OOM。

CASE 7

问题: collect输出大量结果慢

```
/**  
 * Return an array that contains all of the elements in this RDD.  
 */  
def collect(): Array[T] = {  
    val results = sc.runJob(this, (iter: Iterator[T]) => iter.toArray)  
    Array.concat(results: _*)  
}
```

解决:

不要用collect, 直接输出到分布式文件系统。

CASE 8

问题：序列化不给力！

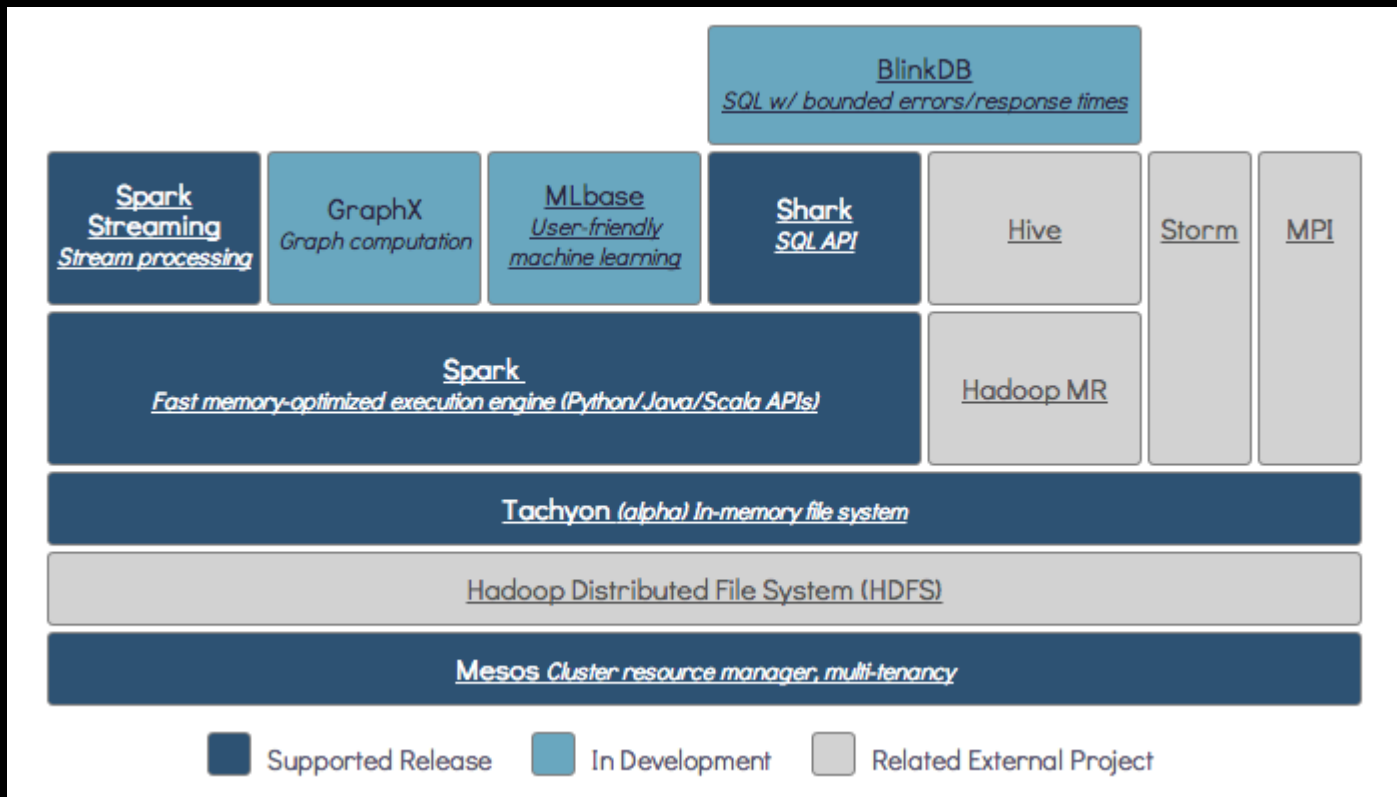
spark默认的序列化机制是JAVA的ObjectOutputStream（兼容性好。但！又大又慢）

解决：

使用Kyro serialization(兼容性略差。但！又快又小)

SPARK就介绍到这，但，这其实只是个开始！

BDAS



Q & A

推荐关注的新浪微博

@hashjoin

@李浩源HY

@明风Andy

@Andrew-Xia

@连城404

@吴甘沙

@尹绪森

还有谁我漏掉了？

thanks all, enjoy spark!

@CrazyJvm